



Supervised classification

- $(x_i)_{i=1,\dots,d}$ features
- $y_i \in \{0, 1\}$ associated labels

Learn a relationship between x and y that allows the model to predict the label y for new unseen data points x .

Supervised classification applications

1. Predictive maintenance of energy equipment

Application: Predicting whether a piece of equipment (e.g. a wind turbine) is likely to break ($y = 1$) or not ($y = 0$) based on measured parameters such as temperature, pressure...

Interpretation:

x : Feature vector representing measurements from the equipment (e.g., temperature, pressure, etc.).

w: Weights indicating the relative importance of each feature in predicting equipment failure.

y: Binary indicator of failure (1 for failure, 0 for normal operation).

2. Classification of buildings by energy performance

Application: Classifying buildings based on their energy efficiency (e.g., low-energy consumption buildings $y = 0$ versus energy-intensive buildings $y = 1$). Interpretation:

x : Features describing the building (e.g., thermal insulation, heating type, surface area, etc.).

w : Contributions of each feature to the likelihood of a building being energy-intensive.

y: Energy performance classification (0 for efficient, 1 for energy-intensive).

◀ ◻ ▶ ◀ ◻ ▶ ◀ ≡ ▶ ◀ ≡ ▶ ≡ 🔍 ↺

Application example

1. Predictive maintenance of energy equipment

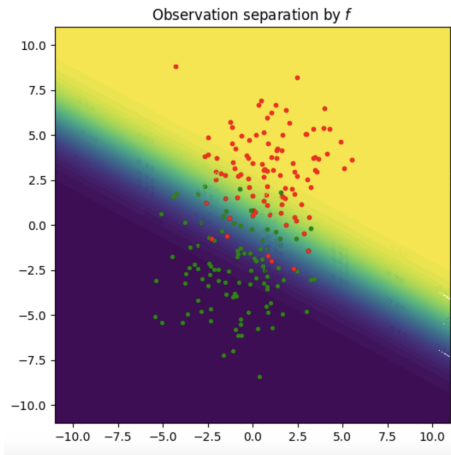
Application: Predicting whether a piece of equipment (e.g. a wind turbine) is likely to break ($y = 1$) or not ($y = 0$) based on measured parameters such as temperature, pressure...

Interpretation:

x : Feature vector representing measurements from the equipment (e.g., temperature, pressure, etc.).

w : Weights indicating the relative importance of each feature in predicting equipment failure.

y : Binary indicator of failure (1 for failure, 0 for normal operation).



Binary supervised classification

Goal

Find the separation line thanks to the training data.

In other words, find the optimal weights

$w = (w_1, w_2) \in \mathbb{R}^2$ and $b \in \mathbb{R}$ such that


$$\langle w, x \rangle + b = w^T x + b = 0.$$

Interpretation:

x : Feature vector representing measurements from the equipment (e.g., temperature, pressure, etc.).

w : Weights indicating the relative importance of each feature in predicting equipment failure.

y : Binary indicator of failure (1 for failure, 0 for normal operation).



Binary supervised classification

How does it work?

Binary supervised classification

Binary supervised classification \rightarrow logistic regression

regression = find a correlation between a binary variable and some observations thanks to an optimization problem

- Decision Trees
- K-Nearest Neighbors (k-NN)
- Probabilistic Models
- Neural Networks


Logistic regression

$$\dots \rightarrow \dots$$

The sigmoid function σ is often used (for f) in logistic regression:

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$

- if $\sigma(\langle w, x \rangle) > 0.5, y = 1$
- if $\sigma(\langle w, x \rangle) < 0.5, y = 0$
- if $\langle w, x \rangle \gg 0, P(y = 1|x) \simeq 1$
- if $\langle w, x \rangle \ll 0, P(y = 1|x) \simeq 0$



Binary supervised classification

Logistic regression

$$x_1 \rightarrow f(\langle w_1, x_1 \rangle)$$

$$\dots \rightarrow \dots$$

$$x_n \rightarrow f(\langle w_n, x_n \rangle)$$


The sigmoid function σ is often used (for f) in logistic regression:

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$

- Likelihood function:

$$\log \left(\frac{P(y=1|x)}{P(y=0|x)} \right) = \langle w, x \rangle$$

- $P(y = 1|x) = \sigma(\langle w, x \rangle)$



Likelihood function

$$P(y = 1|x) = \frac{1}{1 + e^{-\langle w, x \rangle}} := \sigma(\langle w, x \rangle)$$

$$P(y = 0|x) = 1 - \frac{1}{1 + e^{-\langle w, x \rangle}} := 1 - \sigma(\langle w, x \rangle)$$

Log-loss function:

$$f(w) = -\frac{1}{n} \sum_{i=1}^n (y_i \log(\sigma(\langle w, x_i \rangle)) + (1 - y_i) \log(1 - \sigma(\langle w, x_i \rangle))) + \lambda \frac{1}{2} \|w\|^2$$

- y_i : true label (0 or 1),
- $\sigma(\langle w, x_i \rangle)$: probability predicted by the model to get $y_i = 1$.

Likelihood optimization

MINIMIZE the log-loss function:

$$f(w) = -\frac{1}{n} \sum_{i=1}^n (y_i \log(\sigma(\langle w, x_i \rangle)) + (1 - y_i) \log(1 - \sigma(\langle w, x_i \rangle))) + \lambda \frac{1}{2} \|w\|^2$$

→ Gradient descent algorithm !!!

Likelihood optimization

MINIMIZE the log-loss function:

$$f(w) = -\frac{1}{n} \sum_{i=1}^n (y_i \log(\sigma(\langle w, x_i \rangle)) + (1 - y_i) \log(1 - \sigma(\langle w, x_i \rangle))) + \lambda \frac{1}{2} \|w\|^2$$

Remarks:

$\|w\|$ too high is not good:

- if $\langle w, x \rangle \gg 0$, $P(y = 1|x) \simeq 1$
- if $\langle w, x \rangle \ll 0$, $P(y = 1|x) \simeq 0$

Indeed, if $\|w\|$ is high:

- then $\sigma(\langle w, x \rangle)$ becomes too close to 0 or 1, which entails numerical instability problems (NaN).
- then ill-conditioning (problems with the hessian)
- If $\sigma(\langle w, x \rangle) \simeq 0$ or 1: increase too much confidence in our model.

Likelihood optimization

$$\min_w - \frac{1}{n} \sum_{i=1}^n (y_i \log(\sigma(\langle w, x_i \rangle)) + (1 - y_i) \log(1 - \sigma(\langle w, x_i \rangle)))$$

such that $0 \leq \sigma(\langle w, x_i \rangle) \leq 1$ with a penalization for values where $\sigma(\langle w, x_i \rangle) = 0$ or 1 . Instead of the L^2 regularization, we can use the

Interior Point Method (IPM), often used in energy application and find the minimum of:

$$f(w) = -\frac{1}{n} \sum_{i=1}^n (y_i \log(\sigma(\langle w, x_i \rangle)) + (1 - y_i) \log(1 - \sigma(\langle w, x_i \rangle))) + C \sum_{j=1}^d \log(1 + |w_j|^2),$$

where C is the penalization parameter.

→ Gradient descent algorithm !!!

- 1 Interior point method
 - Nonnegative variables

Interior Point Method (IPM)

IPM is a nonlinear optimization algorithm that is particularly effective for solving **large-scale constrained problems**.

It is widely used in energy system optimization with numerous constraints, e.g.:

- Generator production limits
- Electricity transmission constraints (physical laws of the grid)
- Environmental constraints (e.g., CO₂ emissions)

IPM naturally handles these constraints by preventing them from becoming active too early (thanks to the logarithmic barrier function).

Nonnegative variables

Optimization problem. Consider

$$\inf_{x \in \mathbb{R}^n} f(x), \quad \text{subject to: } x_i \geq 0, \quad \forall i = 1, \dots, n.$$

Barrier function: given $c > 0$, consider the function B_c defined by

$$B_c(x) = f(x) - c \sum_{i=1}^n \log(x_i),$$

for all $x \in \mathbb{R}_{>0}^n := \{y \in \mathbb{R}^n \mid y_i > 0, \quad \forall i = 1, \dots, n\}$.

Main idea: approximate (P) by

$$\inf_{x \in \mathbb{R}_{>0}^n} B_c(x). \quad (P_c)$$

Nonnegative variables

General comments.

- We have: $-\log(x_i) \rightarrow \infty$ as $x_i \rightarrow 0$.
 → Feasible points close to the boundary of the feasible set are **penalized** (whatever the value of c).
- A strong modification of the cost function on the feasible set is undesirable.
 → The barrier parameter c should be ideally **very small**.
- Problem (P_c) can be solved with methods for **unconstrained optimization**.
 The standard stepsize rules (Armijo,...) prevents us from getting too close to the boundary.
Ill-conditioning for small values of c .

Nonnegative variables

Example 1.

Consider

$$\inf_{x \in \mathbb{R}} x, \quad \text{subject to: } x \geq 0.$$

- Solution: $\bar{x} = 0$.
- Barrier function: $B_c(x) = x - c \log(x)$.

$$\nabla B_c(x) = 1 - \frac{c}{x} = 0 \iff x = c.$$

Since B_c is convex, $x_c := c$ is the global solution to (P_c) .

- We have $x_c \xrightarrow{c \rightarrow 0} 0$.

Nonnegative variables




Figure: Level-sets $B_c(\cdot)$, for various values of c

Nonnegative variables

Example 2.

Consider

$$\inf_{(x,y) \in \mathbb{R}^2} \frac{1}{2}(y-1)^2 + x, \quad \text{subject to: } \begin{cases} x \geq 0 \\ y \geq 0. \end{cases}$$

- Solution: $(\bar{x}, \bar{y}) = (0, 1)$.
- Solution to the barrier problem: $(x_c, y_c) = \left(c, \frac{1 + \sqrt{1 + 4c}}{2}\right)$.

Nonnegative variables




Figure: Level-sets of $f(\cdot)$

Nonnegative variables




Figure: Level-sets of f , for $c = 0.1$

Nonnegative variables




Figure: Level-sets of $B_c(\cdot)$, for $c = 0.2$

Nonnegative variables




Figure: Level-sets of $B_c(\cdot)$, for $c = 0.5$

Nonnegative variables




Figure: Level-sets of $B_c(\cdot)$, for $c = 1$

Nonnegative variable

Interpretation with the KKT conditions.

Let \bar{x} be a solution to (P) . Let $\bar{\lambda} \in \mathbb{R}^n$ be the associated Lagrange multiplier.

- Lagrangian: $L(x, \lambda) = f(x) - \langle \lambda, x \rangle$. Stationarity condition:

$$\nabla_x L(\bar{x}, \bar{\lambda}) = \nabla f(\bar{x}) - \bar{\lambda} = 0.$$

- Sign condition: $\bar{\lambda}_i \geq 0$.
- Complementarity condition: $\bar{x}_i > 0 \implies \bar{\lambda}_i = 0$.
Equivalently: $\bar{x}_i \bar{\lambda}_i = 0$.

Nonnegative variable

Optimality conditions for the barrier problem.

- For any $x \in \mathbb{R}_{>0}^n$ we denote $\frac{1}{x} = \left(\frac{1}{x_1}, \dots, \frac{1}{x_n}\right)$.
- Let x_c be a solution to (P_c) . We have

$$\frac{\partial B_c}{\partial x_i}(x_c) = \frac{\partial f}{\partial x_i}(x_c) - \frac{c}{x_i} = 0.$$

Therefore $\nabla B_c(x_c) = \nabla f(x_c) - \frac{c}{x_c} = \nabla_x L(x_c, \frac{c}{x_c})$.

- Define $\lambda_c = \frac{c}{x_c} \in \mathbb{R}_{>0}^n$.
The pair (x_c, λ_c) satisfies the KKT conditions approximately:

$$\nabla L(x_c, \lambda_c) = 0, \quad x_{c,i} \lambda_{c,i} = c, \quad \forall i \in \{1, \dots, n\}.$$

Nonnegative variables





Figure: Regularization of the complementarity condition


IPM

$$f(w) = -\frac{1}{n} \sum_{i=1}^n (y_i \log(\sigma(\langle w, x_i \rangle)) + (1 - y_i) \log(1 - \sigma(\langle w, x_i \rangle))) + C \sum_{j=1}^d \log(1 + |w_j|^2),$$

Exercise 5


Code the IPM function and compare the level-set obtained with the functions f without penalization and with the L^2 penalization.

Multilayer perceptron



Multilayer perceptron

One hidden layer



TYPES OF NEURAL NETWORKS

Deep-Feed Forward (DFF)

Support Vector Machine (SVM)

Boltzmann Machine (BM)

Deep Convolutional Network (DCN)

Deconvolutional Network (DN)

I Neural Turing Machine (NTM)

BFGS

Quasi-Newton

$$H_k = H_{k-1} + \frac{(y_{k-1} - H_{k-1} d_{k-1}) d_{k-1}^T}{d_{k-1}^T d_{k-1}}$$

with $d_{k-1} = x_k - x_{k-1}$ and $y_{k-1} = \nabla f(x_k) - \nabla f(x_{k-1})$

BFGS

Quasi-Newton

$$H_k^{-1} = \left(I - \frac{\bar{d}_{k-1} y_{k-1}^T}{\bar{d}_{k-1}^T y_{k-1}} \right) H_{k-1}^{-1} \left(I - \frac{\bar{d}_{k-1} y_{k-1}^T}{\bar{d}_{k-1}^T y_{k-1}} \right) + \frac{\bar{d}_{k-1} \bar{d}_{k-1}^T}{\bar{d}_{k-1}^T y_{k-1}}$$

with $\bar{d}_{k-1} = x_k - x_{k-1}$ and $y_{k-1} = \nabla f(x_k) - \nabla f(x_{k-1})$

$$x_k = x_{k-1} + \alpha d$$

$$\bar{d} = \alpha d$$

- Update d
- Update α thanks to Armijo/ Wolfe
- Update x
- Update y, \bar{d}
- Find H^{-1}